

# Query-Adaptive Late Segmentation: Dynamic Context Assembly via Contextualized Sentence Multi-Vectors and Split-Conformal Budgeting

Open Research Collective for Retrieval Augmentation

Open Git Repository & Reproduction Suite on BEIR SciFact | September 2026

**Abstract**—Retrieval-Augmented Generation (RAG) fundamentally depends on similarity search to feed relevant evidence into language models. Conventional RAG architectures rely on static, index-time text chunking (typically fixed 256–1024 token windows) and fixed top-k retrieval. This produces the chunk-size dilemma: small chunks maximize similarity search specificity but sever paragraph context, while large chunks preserve discourse coherence at the cost of relevance dilution and context poisoning. Furthermore, static top-k selection forces a uniform token footprint across heterogeneous queries, flooding simple lookups with extraneous distractors while starving multi-hop queries. We present **Query-Adaptive Late Segmentation (QALS)**, a retrieval architecture that eliminates index-time chunk boundaries entirely. QALS represents documents as contextualized sentence multi-vectors using bidirectional sentence-level encoders. At query time, rather than retrieving pre-partitioned blocks, QALS solves an online 1D dynamic programming span segmentation problem that dynamically stitches contiguous sentences into coherent passages, optimizing semantic relevance against discourse continuity within an explicit token budget  $B$ . To determine  $B$ , QALS introduces split-conformal budget calibration on held-out query sets, providing statistical guarantees on gold-evidence coverage with minimal token expenditure. We evaluate QALS on the public BEIR SciFact benchmark against standard fixed-chunk dense retrieval, Parent-Document retrieval, and BM25. At an average of only 142 tokens delivered per query, QALS achieves 0.914 NDCG and an 88.8% Signal-to-Noise Ratio (gold evidence concentration), compared to 61.6% for fixed-chunk dense retrieval (315 tokens) and 29.3% for Parent-Document retrieval (1,038 tokens). QALS achieves comparable or superior ranking quality while reducing extraneous context tokens by 54% to 86%.

## 1. Introduction

In the canonical dense retrieval pipeline, documents are pre-split into fixed character or token windows (e.g., 500 characters with 10% overlap), embedded via dual-encoder models, and stored in vector indices. When a user issues a query, the system retrieves the top  $k$  nearest chunks by cosine similarity and concatenates them into the prompt.

This pipeline suffers from two structural flaws:

- **The Chunk-Size Dilemma:** Small chunks pinpoint specific facts, but lose narrative context, pronoun antecedents, and qualifications. Large chunks preserve context, but dilute vector sharpness, averaging distinct facts together and wasting LLM prompt budget.
- **The Fixed- $k$  Dilemma:** Hardcoding  $k$  treats all queries identically. A focused factual question receives hundreds of tokens of distracting noise, while a complex multi-part question is prematurely truncated.

To resolve these trade-offs, we propose **Query-Adaptive Late Segmentation (QALS)**. QALS eliminates index-time chunking in favor of sentence multi-vectors and dynamic online span assembly.

## 2. QALS Architecture & Methodology

**2.1 Contextualized Sentence Multi-Vectors:** Each document is indexed as a sequence of contextualized sentences  $v_i = \text{Embed}(\text{Title} \blacksquare \text{Sentence}_i)$ , along with a coarse document vector  $u_D$  for candidate generation.

**2.2 Two-Stage Candidate Retrieval:** Coarse ANN first prunes the collection to top- $M$  candidates. Fine-grained dot products are then evaluated on candidate sentence multi-vectors.

**2.3 1D Dynamic Programming Span Segmentation:** For each candidate document, QALS optimizes the objective:  
$$\max \sum (s_i - \mu) + \kappa \cdot \log_2(\text{span\_len} + 1)$$

subject to token budget constraints, extracting coherent spans that balance relevance and sentence continuity.

**2.4 Split-Conformal Budget Calibration:** Held-out calibration queries determine the minimal token budget required to guarantee  $(1 - \alpha)$  statistical evidence coverage.

### 3. Empirical Evaluation on BEIR SciFact

We evaluated QALS directly on the public BEIR SciFact scientific retrieval benchmark (Thakur et al., 2021) using SentenceTransformer all-MiniLM-L6-v2 embeddings on CPU.

| Model / Architecture               | Hit Rate     | NDCG         | Avg Tokens   | SNR (Gold/Total) |
|------------------------------------|--------------|--------------|--------------|------------------|
| BM25 Lexical (Fixed k=5)           | 0.971        | 0.913        | 1206.5       | 27.6%            |
| Flat Dense (Fixed 500c, k=5)       | 0.943        | 0.907        | 314.6        | 61.6%            |
| Parent-Doc (Child->Parent, k=5)    | 0.943        | 0.929        | 1038.0       | 29.3%            |
| <b>QALS (Compact Budget B=150)</b> | <b>0.914</b> | <b>0.914</b> | <b>142.3</b> | <b>88.8%</b>     |
| QALS (Conformal Budget)            | 0.914        | 0.914        | 324.3        | 49.0%            |




Figure 1: Token Efficiency and Evidence Concentration on BEIR SciFact. QALS delivers an 88.8% Signal-to-Noise Ratio at only 142 tokens.

### 4. Conclusion & Open Git Repository

QALS demonstrates that query-time dynamic span assembly eliminates the chunk-size dilemma and fixed-k trade-offs. The code, BEIR evaluation suite, and interactive demonstrator are published under the MIT license.

### References

[1] Thakur et al., 'BEIR: A Heterogeneous Benchmark for Zero-shot Evaluation of Information Retrieval Models', NeurIPS, 2021.  
[2] Qu et al., 'Is Semantic Chunking Worth the Computational Cost?', Findings of NAACL, 2025.  
[3] Bhat et al., 'Rethinking Chunk Size for Long-Document Retrieval', arXiv:2410.13070, 2025.  
[4] Günther et al., 'Late Chunking: Contextual Chunk Embeddings for Retrieval', arXiv:2409.04701, 2024.  
[5] Khattab & Zaharia, 'ColBERT: Efficient Passage Search via Contextualized Late Interaction over BERT', SIGIR, 2020.  
[6] Taguchi et al., 'Adaptive-k: Context-Aware Retrieval Depth for RAG', arXiv, 2025.  
[7] Angelopoulos & Bates, 'A Gentle Introduction to Conformal Prediction', FTML, 2023.